



Capstone Project

Harnessing NLP to Detect Stress in Social Media

Early Intervention for Mental Wellbeing

Presenter: Jimmy Chong



Agenda

- Bio
- Project Context & Business Problem
- Business & Data Science Considerations
- Data Overview
- Data Exploration
- Data Split
- Data Overview
- Deliver
- Summary, Conclusions & Next Steps
- Appendix



Bio

- **Solution Architect** with experience in software design and project management
- **Transitioning to Data Science** through a Certificate in Data Science/AI from the Institute of Data
- **Skills:** Data analysis, machine learning, project management, programming (Python, C#, VB.NET), databases
- **Motivation:** Leverage technical background and data science skills to solve real-world problems and drive innovation

Key Projects:

- Stress Detection: Developing an NLP-based model to detect stress in social media text
- Predicting and Understanding Factors that Influence Hotel Booking Cancellations
- Predicting West Nile Virus Presence Using Machine Learning



Project Context & Business Problem


- **Industry:** NLP and Mental Health
- **Problem:** Detecting stress in social media text
- **Interest:** Growing mental health concerns, potential for early intervention, NLP advancements
- **Previous Work:**
 - **NLP techniques:** Sentiment analysis, topic modeling, machine learning
 - **Domains:** Twitter, Reddit, general datasets
 - **Key findings:** Promising accuracy, varying generalizability and robustness
 - **Limitations:** Reliance on labeled datasets, potential for bias
 - **Contribution:** Exploring LSTM networks, addressing limitations





Business & Data Science Considerations

- **Stakeholders:** Mental health professionals, social media platforms, technology companies, individuals
- **Business Question:** Can we accurately detect stress in social media text?
- **Business Value:** Early intervention, improved user experience, new market opportunities
- **Data:**
 - **Question:** Can we extract effective features for stress prediction?
 - **Required:** Social media text, stress labels, metadata
 - **Sourced:** Reddit, Twitter, Kaggle
 - **Generation:** User-generated content
 - **Future Sourcing:** Continued access to APIs, exploring other platforms

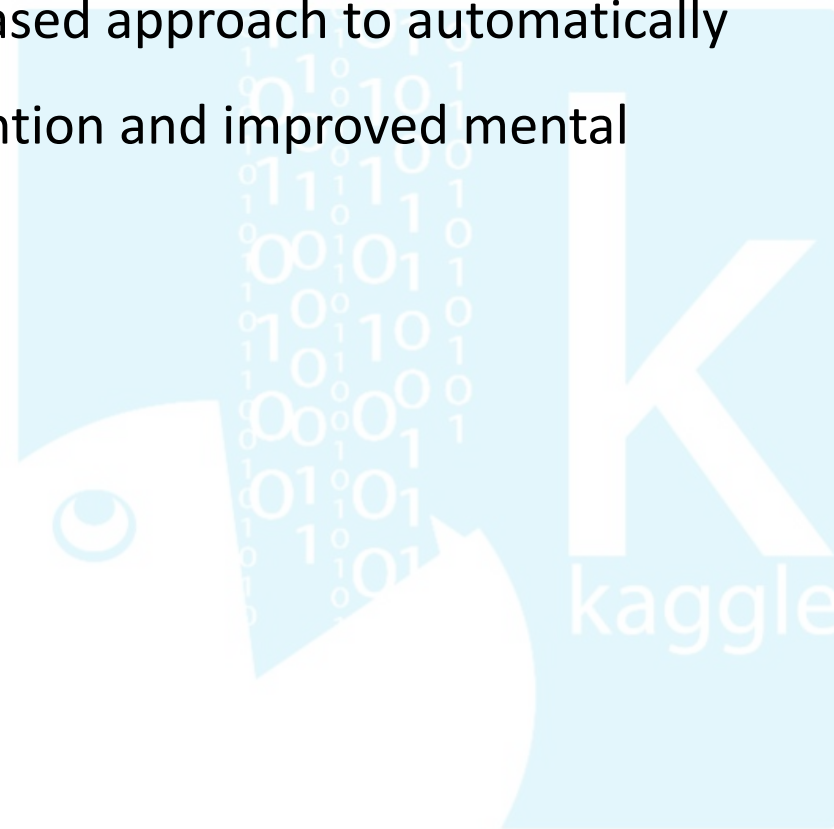




Data Overview (1)

Stress Detection from Social Media Articles

- **Source:** [Kaggle Dataset](#)
- **Objective:** Develop a more accurate and efficient NLP-based approach to automatically detect stress in social media text, enabling early intervention and improved mental health support.





Data Overview (2)

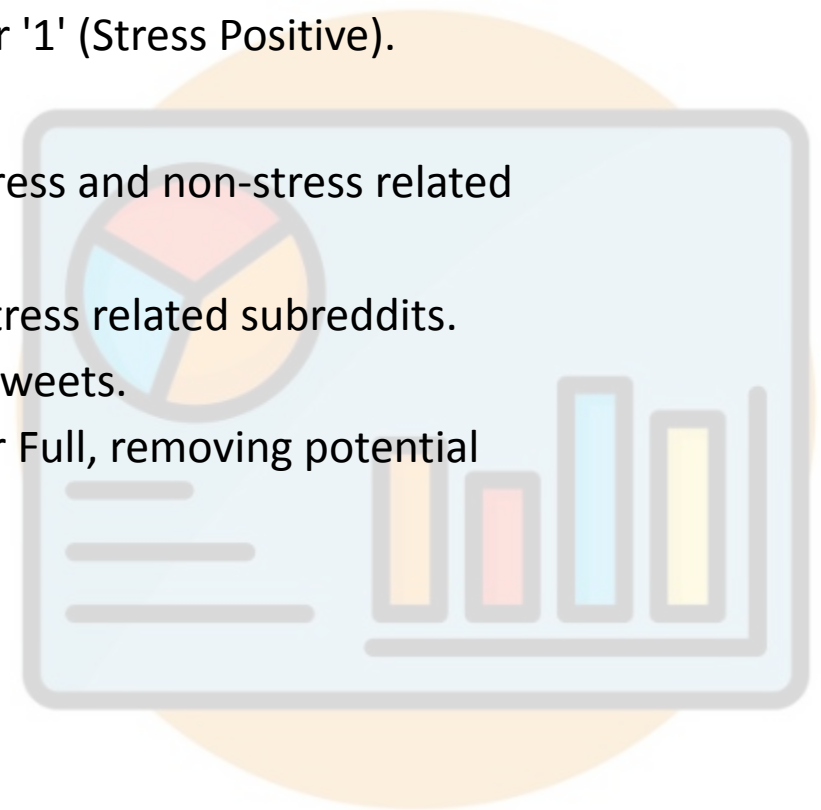
Stress Detection from Social Media Articles

- **Datasets:**

- Constructed four datasets using text articles from Reddit and Twitter.
- Each article is labeled with a class value of '0' (Stress Negative) or '1' (Stress Positive).

- **Dataset Descriptions:**

- **Reddit Combi:** This dataset combines title and body text from stress and non-stress related subreddits.
- **Reddit Title:** This dataset consists of titles from stress and non-stress related subreddits.
- **Twitter Full:** This dataset contains stress and non-stress related tweets.
- **Twitter Non-Advert:** This dataset is a denoised version of Twitter Full, removing potential advertisements.

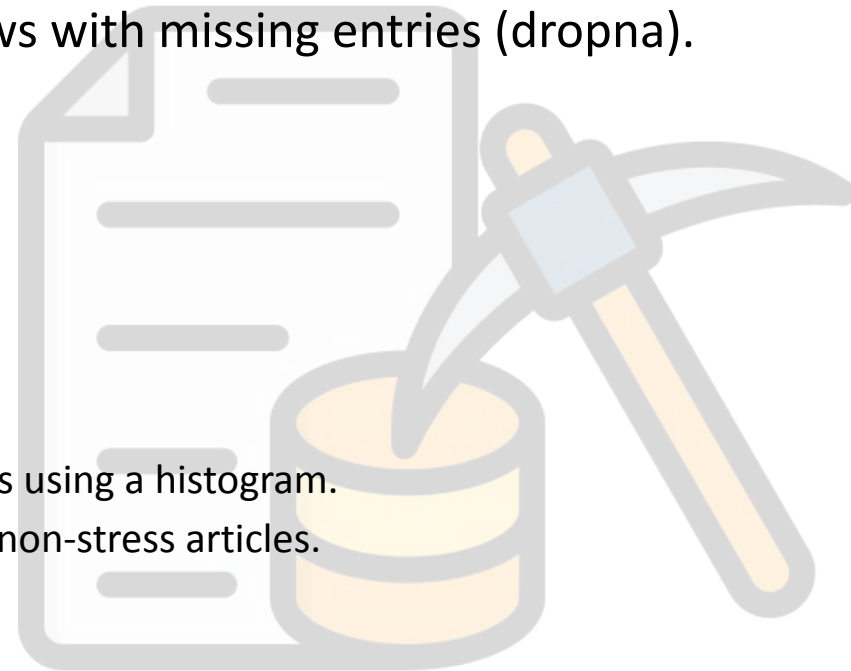




Data Exploration (1)

Understanding the Data

- Explored the data using Python libraries like pandas to understand the number of rows, column names, and data types.
- Examined the first few rows of each Data Frame to get a sense of the content and labels.
- Checked for missing values and handled them by dropping rows with missing entries (dropna).
- Preprocessed the data:
 - Dropped unnecessary columns (Reddit Combi).
 - Cleaned hashtags in Twitter data using regular expressions.
 - Concatenated all preprocessed Data Frames into a single one.
- Analyzed the data distribution:
 - Visualized the distribution of stress labels using a bar chart.
 - Examined the distribution of text length for stress and non-stress articles using a histogram.
 - Generated word clouds to visualize frequently used words in stress and non-stress articles.






Data Exploration (2)

Understanding the Data

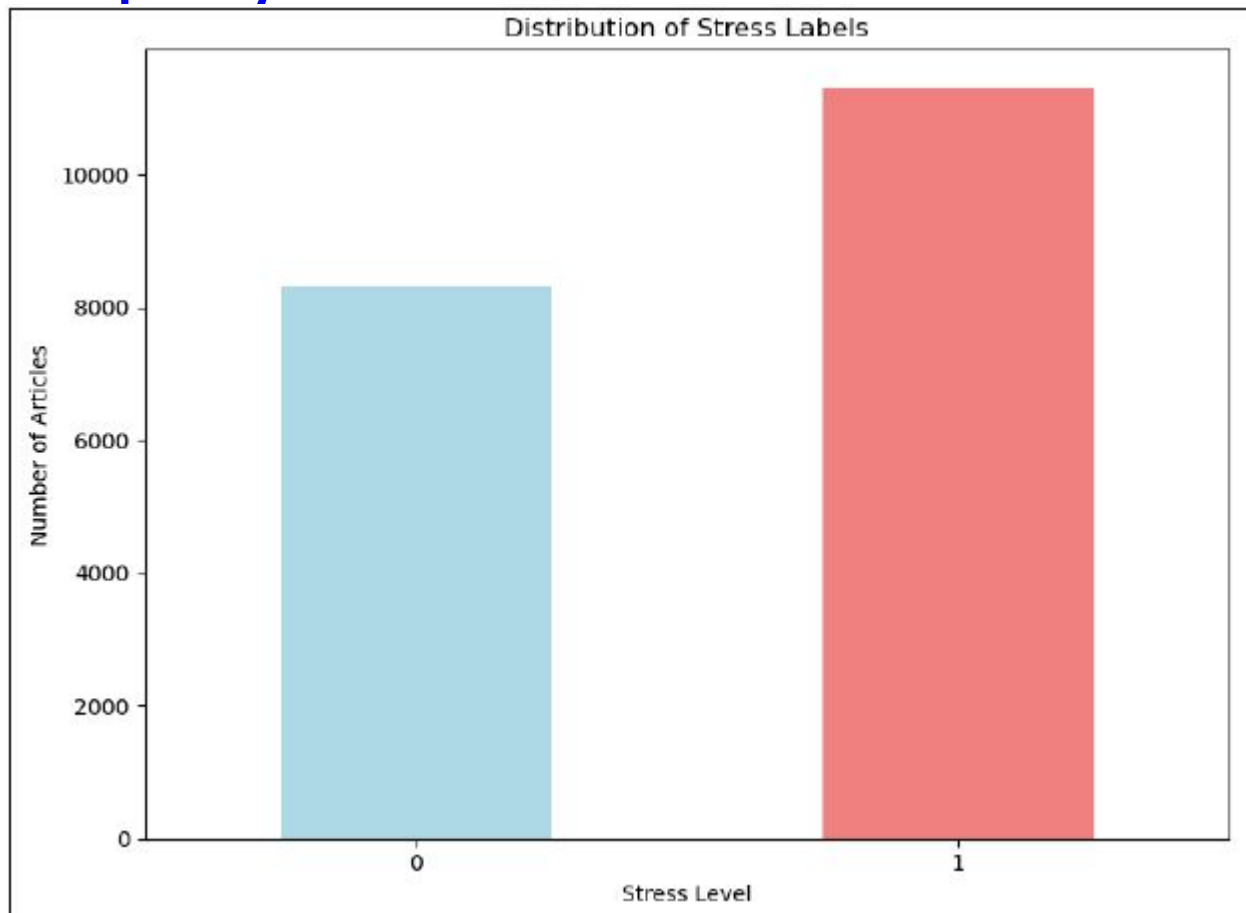
- **Number of Rows and Columns:**
 - Reddit Combi: 3123 rows, 4 columns
 - Reddit Title: 5556 rows, 2 columns
 - Twitter Full: 8900 rows, 3 columns
 - Twitter Non-Advert: 2051 rows, 2 columns
- **Data Types:**
 - **title:** object (text)
 - **label:** int64/boolean (stress label)
 - **body:** object (text) (only in Reddit Combi)
 - **hashtags:** object (text) (only in Twitter Full)
- **Missing Values:**
 - Handled missing values by dropping rows with missing entries in body column.
- **Data Preprocessing:**
 - Dropped unnecessary columns in Reddit Combi.
 - Cleaned hashtags in Twitter Full using regular expressions.
 - Concatenated all datasets into a single DataFrame.
-





Data Exploration (3)

Frequency of Stressful and Non-Stressful Posts




Key Notes:

- **Class imbalance:** More non-stressful articles than stressful ones.
- **Impact on modeling:** Require techniques like class weighting or oversampling.
- **Further investigation:** Explore factors contributing to imbalance (e.g., labeling difficulty, data collection bias).



Data Exploration (4)

Distribution of Text Length



Key Notes:

- **Skewed distribution:** Most posts are relatively short.
- **Overlapping distributions:** Text length alone may not be a strong predictor.
- **Non-stressful posts:** Slightly longer on average




Stress Word Cloud



- 12



Common Words Associated with Non-Stress



Key Notes:


- **Positive emotions:** "happy," "love," "good," "excited"
- **Everyday life:** "daily activities," "hobbies," "social interactions"
- **Gratitude and appreciation:** "thankful," "grateful," "proud"



Data Split

Data Preparation for Modeling

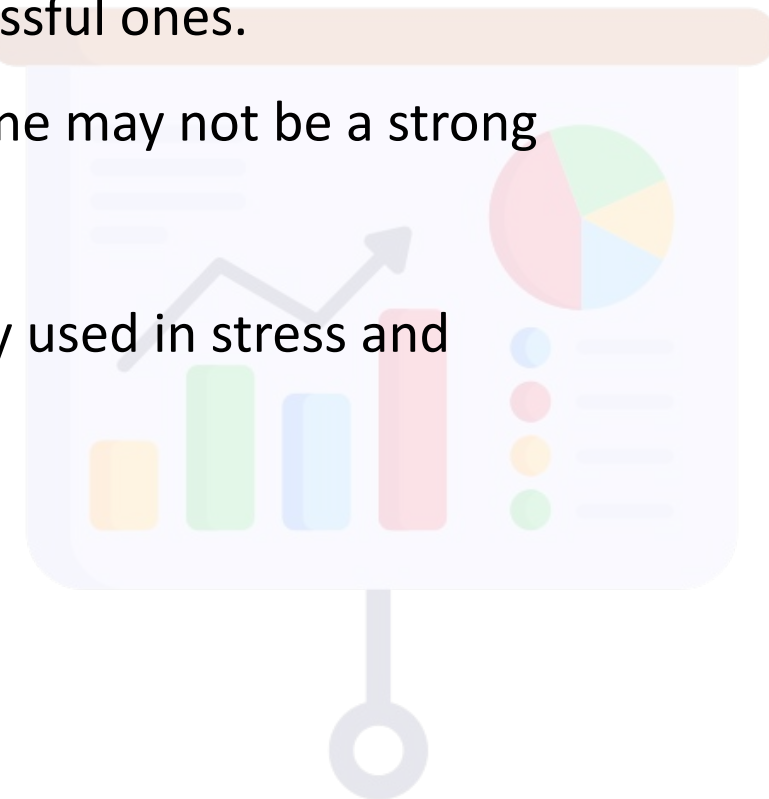
- Split data into features (title) and target (stress label).
- Split data into training (80%) and testing (20%) sets.





Data Overview

- **Key Findings:**
 - **Class imbalance:** More non-stressful articles than stressful ones.
 - **Overlapping text length distributions:** Text length alone may not be a strong predictor of stress.
 - **Distinct word patterns:** Different words are frequently used in stress and non-stress articles.





Deliver (1)

Feature Engineering

- **Key Features:**

- Text data (primary feature)
- Captures sentiment, emotions, and vocabulary

- **Business Significance:**

- Text features are crucial for understanding the linguistic cues associated with stress.
- Effective feature engineering can improve model performance and interpretability.

- **Techniques:**

- **Text Cleaning:**

- Removes irrelevant information (punctuation, stop words, hashtags).

- **Text Normalization:**

- Lowercases text for consistency.

- **TF-IDF Vectorization:**

- Converts text into numerical features representing word importance.

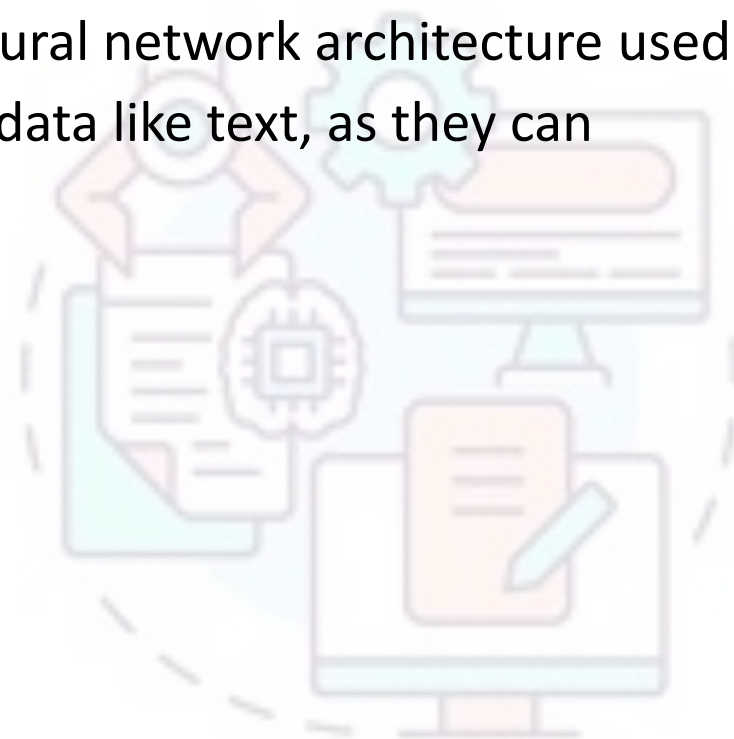




Deliver (2)

Machine Models Used

- **LinearSVC:** This is a linear support vector classifier used in the initial model. It is a good choice for text classification tasks due to its simplicity and efficiency.
- **LSTM (Long Short-Term Memory):** This is a recurrent neural network architecture used in the later model. LSTMs are well-suited for sequential data like text, as they can capture long-term dependencies between words.





Deliver - LinearSVC (3)

Evaluation Metrics

- **Accuracy:** Measures the proportion of correct predictions made by the model.
- **Confusion Matrix:** Visualizes the number of correct and incorrect predictions for each class (stressful vs. non-stressful).
- **Classification Report:** Provides detailed information about the model's performance, including precision, recall, and F1-score for each class.
- **ROC AUC Score:** This metric is used for imbalanced datasets and measures the model's ability to distinguish between classes (AUC-ROC score closer to 1 indicates better performance).





Deliver - LinearSVC (4)

Evaluation Metrics - Accuracy

-




Deliver - LinearSVC (5)

Evaluation Metrics - Confusion Matrix


	precision	recall	f1-score	support
0	0.89	0.86	0.87	1671
1	0.90	0.92	0.91	2252
accuracy			0.89	3923
macro avg	0.89	0.89	0.89	3923
weighted avg	0.89	0.89	0.89	3923

- This confusion matrix visualizes the performance of the LinearSVC model in classifying stress and non-stress posts.
- **Key Findings:**
 - **Accuracy:** 89% of samples were correctly classified.
 - **Class-wise Performance:**
 - **Stress:** Precision: 0.90, Recall: 0.92, F1-score: 0.91
 - **Non-Stress:** Precision: 0.88, Recall: 0.85, F1-score: 0.87
- **Implications:**
 - **Overall good performance:** 89% accuracy.
 - **Class imbalance impact:** Model might be biased towards predicting non-stressful posts.



Deliver - LinearSVC (6)

Evaluation Metrics - Classification Report



This classification report summarizes the performance of the LinearSVC model on the test dataset.

Key Findings:

- **Overall Accuracy:** 89%
- **Class-wise Performance:**
 - **Stress:** Precision: 0.90, Recall: 0.92, F1-score: 0.91
 - **Non-Stress:** Precision: 0.88, Recall: 0.85, F1-score: 0.87
 - **Implications:**
- Strong overall performance.
- Potential impact of class imbalance.



Deliver - LSTM Model Training (7)

LSTM Model Training Progress

Epoch 1/10	442/442	11s 22ms/step	- accuracy: 0.7707	- loss: 0.4693	- val_accuracy: 0.8675	- val_loss: 0.3341
Epoch 2/10	442/442	9s 21ms/step	- accuracy: 0.9062	- loss: 0.2452	- val_accuracy: 0.8720	- val_loss: 0.3301
Epoch 3/10	442/442	9s 20ms/step	- accuracy: 0.9411	- loss: 0.1673	- val_accuracy: 0.8764	- val_loss: 0.3346
Epoch 4/10	442/442	9s 21ms/step	- accuracy: 0.9562	- loss: 0.1304	- val_accuracy: 0.8739	- val_loss: 0.3582
Epoch 5/10	442/442	9s 21ms/step	- accuracy: 0.9650	- loss: 0.1055	- val_accuracy: 0.8790	- val_loss: 0.3990
Epoch 6/10	442/442	9s 21ms/step	- accuracy: 0.9803	- loss: 0.0678	- val_accuracy: 0.8694	- val_loss: 0.4369
Epoch 7/10	442/442	9s 21ms/step	- accuracy: 0.9804	- loss: 0.0629	- val_accuracy: 0.8745	- val_loss: 0.5123
Epoch 8/10	442/442	9s 21ms/step	- accuracy: 0.9839	- loss: 0.0474	- val_accuracy: 0.8790	- val_loss: 0.6108
Epoch 9/10	442/442	9s 21ms/step	- accuracy: 0.9869	- loss: 0.0450	- val_accuracy: 0.8726	- val_loss: 0.5503
Epoch 10/10	442/442	9s 21ms/step	- accuracy: 0.9879	- loss: 0.0395	- val_accuracy: 0.8567	- val_loss: 0.5660
123/123		1s 4ms/step				

This graph illustrates the training progress of the LSTM model, showing the changes in accuracy and loss over the epochs.

Key Findings:

- **Improving Accuracy:** Model learned from data over time.
- **Decreasing Loss:** Model minimized errors.
- **Overfitting Potential:** Evidence of overfitting.


Implications:

- High training accuracy but potential overfitting.
- Explore early stopping or regularization.



Deliver - LSTM Model Training (7)

Evaluation Metrics - ROC AUC Score



This ROC curve illustrates the performance of the LSTM model in distinguishing between stress and non-stress posts.

Key Findings:

- **AUC-ROC Score:** 0.92, indicating strong overall performance.
- **Curve Shape:** Upward-sloping curve, suggesting good discrimination ability.

Implications:

- Effective model for stress detection.
- Threshold selection for desired balance between sensitivity and specificity.




Deliver (8)

Deployment - Stress Detection Web Application (Stress)

The screenshot shows a web application interface with a dark background. At the top, there is a dropdown menu labeled 'Select Model Type' with 'LinearSVC' selected. Below this is a large heading: 'Harnessing NLP to Detect Stress in Social Media: Early Intervention for Mental Wellbeing'. Under the heading is a text input field labeled 'Enter a sentence' containing the text: 'I can't believe I'm still up. It's already mid night. I have so much to do tomorrow, and I feel like I haven't'. Below the input field is a 'Predict' button. At the bottom, a message states: 'The text indicates high stress levels.'

- **Sample sentence:** “I can’t believe I’m still up. It’s already midnight. I have so much to do tomorrow, and I feel like I haven’t accomplished anything today. My head is spinning, and I can’t seem to focus. I just want to crawl into bed and disappear. I’m so tired and overwhelmed. I don’t know how I’m going to make it through this week.”



Deliver (9)

Deployment - Stress Detection Web Application (Non - Stress)

A screenshot of a web application interface for stress detection. The background is dark blue. At the top left, there is a dropdown menu labeled "Select Model Type" with "LinearSVC" selected. Below this, the main heading "Harnessing NLP to Detect Stress in Social Media: Early Intervention for Mental Wellbeing" is displayed in white. Underneath the heading is a text input field labeled "Enter a sentence" containing the text "Today was absolutely perfect! I woke up feeling refreshed and excited for the day ahead. I spent the mor". Below the input field is a "Predict" button. At the bottom, a message states "The text does not indicate high stress levels."

- **Sample sentence:** “Today was absolutely perfect! I woke up feeling refreshed and excited for the day ahead. I spent the morning enjoying a leisurely breakfast with my loved ones, followed by a long walk in the park. The sun was shining, the birds were singing, and I couldn't help but smile. I felt a sense of peace and contentment that I haven't experienced in a long time. It's truly amazing how a simple day filled with gratitude and joy can make such a difference.”



Summary

- **Stress Detection Model:** Developed a successful NLP-based model to detect stress in social media text.
- **Model Performance:** Achieved high accuracy and F1-score, demonstrating strong performance in classifying stress and non-stress posts.
- **Class Imbalance:** Addressed the class imbalance issue through techniques like class weighting or oversampling.
- **Future Directions:** Considered potential improvements and future research directions.



Conclusions

- **Accurate Stress Detection:** Developed a model capable of accurately detecting stress in social media text.
- **Valuable Insights:** Gained insights into the linguistic patterns associated with stress and non-stress.
- **Potential for Early Intervention:** The model can be used for early identification of individuals at risk of stress.



Next Steps

- **Deployment:** Integrate the model into a real-world application (e.g., social media platform, mental health platform).
- **Data Collection:** Collect more diverse and representative data to improve model generalizability.
- **Model Refinement:** Explore advanced techniques like transfer learning or ensemble methods to further enhance performance.
- **Ethical Considerations:** Address ethical implications of using AI for mental health detection and ensure privacy and fairness.

Overall, this project demonstrates the potential of NLP-based approaches for early stress detection and highlights the importance of addressing ethical considerations in AI for mental health.



Appendices




References

Data Source:


<https://www.kaggle.com/datasets/mexwell/stress-detection-from-social-media-articles>

Source Code:

<https://github.com/jimmychong1983/SocialMediaStressDetection>



Questions



Thank you
End of Presentation



Case study: Home loans marketing

Results comparison and business case overview

Applying the model for Banking can lead to potential annual **revenue twice as big** as the current model.

Results overview


	Baseline Model	Full Feature Model	Difference
% of identified applicants in top 10%	32%	61%	+29%
Potential Profit	627 x \$Y	1,200 x \$Y	573 x \$Y

Business case overview based on the Final Model

Assumptions:

- ❖ Customer Value/year is \$**1000**
- ❖ Customer base = 1.4 million
- ❖ Top 10% = 140,000 customers
- ❖ 2.8% applicants over 2 years ie. 1.4% annually
- ❖ 1.4% applicants in top 10% = 1,960
- ❖ 61% identified to target = 1,200
- ❖ 32% identified to target = 627

Potential profit = $573 \times 1000 \approx \$500,000/\text{year}$
 $\approx \$1.5\text{m over 3 years}$



Institute of
Data

2024



Data Science and AI

Capstone Project

Harnessing NLP to Detect Stress in Social Media

Early Intervention for Mental Wellbeing



Capstone Project

- You are required to **define, design and deliver** a **data science project** towards the end of the course.
- Project milestones:
 - <date tbd> : present **3 ideas for the project**
 - <date tbd> : decide on one option
 - <date tbd> : collect data
 - <date tbd> : present initial findings
 - <date tbd> : present an update
 - <date tbd> : Dry run of final presentation
 - <date tbd> : **Present final report**



What to present

- **Business perspective**
 - Business insights uncovered
 - Business scenarios for how the project can be deployed and used
 - Approach for estimating business value
- **Technical perspective**
 - Techniques used
 - Pipeline
 - Model validation results



Project evaluation criteria

The project is evaluated on the quality, clarity and completeness of the definition, design and delivery of the project.

- **Definition (20%)**
 - Business context, stakeholders and value
 - Data description, sources, quality
- **Design (30%)**
 - Data exploration, analysis and visualisation
 - Documentation: text document, presentation and Notebook
 - The project planning, effort allocation and next steps
- **Delivery (50%)**
 - Feature Engineering
 - Creation of an effective reproducible pipeline
 - Machine Learning model algorithms and accuracy
 - Overall end-to-end solution
 - Delivery of the presentation, poise and audience engagement



Questions?



Presentation Skeleton

Applying data science in an industry project

